

```

!pip install python-mnist
from __future__ import print_function
import tensorflow as tf
import keras
from keras.layers import Dense
from keras.datasets import mnist
from keras.models import Sequential, Model, load_model
from keras.layers import Dense, Dropout, Activation, Flatten, Input
from keras.layers import Conv2D, MaxPooling2D, UpSampling2D, Reshape
from keras.layers.advanced_activations import LeakyReLU
from sklearn.metrics import classification_report
from sklearn.model_selection import train_test_split
import time, torch
from keras.datasets import mnist
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

```

▼ Data Loading & Preprocessing

```

def data_prepare():
    """
    this function loads the mnist data, normalizes pixel values, converts labels to one hot enc
    """

    (x_train, y_train), (x_test, y_test) = mnist.load_data()
    print("Training data", x_train.shape, "Test data:", x_test.shape)

    x_train = x_train.astype('float32')
    x_test = x_test.astype('float32')

    x_train = x_train.reshape(x_train.shape[0], x_train.shape[1], x_train.shape[2], 1)
    x_test = x_test.reshape(x_test.shape[0], x_test.shape[1], x_test.shape[2], 1)

    # Normalization b/w 0 and 1
    x_train = (x_train/127.5) - 1
    x_test = (x_test/127.5) - 1

    # converting output to categorical vector of dim = num_classes
    y_train = keras.utils.to_categorical(y_train)
    y_test = keras.utils.to_categorical(y_test)

    return x_train, y_train, x_test, y_test

```


```

x_train, y_train, x_test, y_test = data_prepare()

```

Training data (60000, 28, 28) Test data: (10000, 28, 28)

Input shape for CNN of x & y of training (60000, 28, 28, 1) (60000, 10) & test data (10



```
batch_size = 128
num_classes = 10
epochs = 50
img_shape = (x_train.shape[1], x_train.shape[2], 1)

#This is the input dimension for gan
latent_dim = 50
```

▼ Setting up Generator

```
def generator():
    """
    This function returns a generator model framework based on CNN
    """
    model_gen = Sequential()

    # An input of 50 dimension is given to the Dense layer, the weights are initialized with 'truncated_normal'
    model_gen.add(Dense(7*7*128, input_dim = latent_dim, kernel_initializer = tf.keras.initializers.truncated_normal))
    #Leaky ReLU addresses the problem of dead gradients with ReLU and therefore is used
    model_gen.add(LeakyReLU(alpha=0.2))
    model_gen.add(Reshape((7, 7, 128)))

    #Doubling the x and y size
    model_gen.add(UpSampling2D())
    #He normal initialisation is observed to work well with Leaky ReLU activations
    model_gen.add(Conv2D(128, kernel_size=(3,3), kernel_initializer= tf.keras.initializers.he_normal))
    #Leaky ReLU is found to be a best practice for GAN
    model_gen.add(LeakyReLU(alpha=0.2))
    model_gen.add(UpSampling2D())
    model_gen.add(Conv2D(64, (3, 3), kernel_initializer= tf.keras.initializers.he_normal(seed = 1)))
    model_gen.add(LeakyReLU(alpha=0.2))

    #Gives output of size 28*28 with one channel
    model_gen.add(Conv2D(1, kernel_size=(3,3), padding="same"))

    #makes elementwise output between -1 to 1
    model_gen.add(Activation("tanh"))
    #Adam optimizer is found to work well for generator
    model_gen.compile(loss = 'binary_crossentropy', optimizer = keras.optimizers.Adam(learning_rate=0.0002))
    return model_gen

g = generator()
g.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 6272)	319872
leaky_re_lu (LeakyReLU)	(None, 6272)	0
reshape (Reshape)	(None, 7, 7, 128)	0
up_sampling2d (UpSampling2D)	(None, 14, 14, 128)	0
conv2d (Conv2D)	(None, 14, 14, 128)	147584
leaky_re_lu_1 (LeakyReLU)	(None, 14, 14, 128)	0
up_sampling2d_1 (UpSampling2D)	(None, 28, 28, 128)	0
conv2d_1 (Conv2D)	(None, 28, 28, 64)	73792
leaky_re_lu_2 (LeakyReLU)	(None, 28, 28, 64)	0
conv2d_2 (Conv2D)	(None, 28, 28, 1)	577
activation (Activation)	(None, 28, 28, 1)	0
Total params: 541,825		
Trainable params: 541,825		
Non-trainable params: 0		

▼ Setting up Discriminator

```
def discriminator():  
    """  
    This function returns a discriminator classification model framework based on CNN  
    """  
    model_dis = Sequential()  
  
    #64 kernels with size (3,3) will be employed, zero padding is kept so output size same as i  
    model_dis.add(Conv2D(128, kernel_size= (3, 3), strides = (2, 2), padding = 'same', input_s  
    model_dis.add(LeakyReLU(alpha=0.2))  
    #dropout helps avoid overfitting  
    model_dis.add(Dropout(0.3))  
  
    model_dis.add(Conv2D(64, kernel_size= (3, 3), strides = (2, 2), padding = 'same'))  
    model_dis.add(LeakyReLU(alpha=0.2))  
    model_dis.add(Dropout(0.3))  
  
    model_dis.add(Conv2D(32, kernel_size= (3, 3), strides = (2, 2), padding = 'same'))  
    model_dis.add(LeakyReLU(alpha=0.2))  
    model_dis.add(Dropout(0.3))
```

```

#output is converted into 1d tensor by serially arranging all elements
model_dis.add(Flatten())

#sigmoid ensures output to be between 0 and 1
model_dis.add(Dense(1, activation = 'sigmoid'))
model_dis.compile(loss = "binary_crossentropy", optimizer = keras.optimizers.Adam(learning_

return model_dis

d = discriminator()
d.summary()

```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
=====		
conv2d_3 (Conv2D)	(None, 14, 14, 128)	1280
leaky_re_lu_3 (LeakyReLU)	(None, 14, 14, 128)	0
dropout (Dropout)	(None, 14, 14, 128)	0
conv2d_4 (Conv2D)	(None, 7, 7, 64)	73792
leaky_re_lu_4 (LeakyReLU)	(None, 7, 7, 64)	0
dropout_1 (Dropout)	(None, 7, 7, 64)	0
conv2d_5 (Conv2D)	(None, 4, 4, 32)	18464
leaky_re_lu_5 (LeakyReLU)	(None, 4, 4, 32)	0
dropout_2 (Dropout)	(None, 4, 4, 32)	0
flatten (Flatten)	(None, 512)	0
dense_1 (Dense)	(None, 1)	513
=====		
Total params: 94,049		
Trainable params: 94,049		
Non-trainable params: 0		

▼ Combining Generator and Discriminator

```

def gan(generator, discriminator):
    """
    This function integrates the generator and framework into a sequence and then returns a moc
    """
    # We will train generator and discriminator sequentially and not parallely. We train discri
    discriminator.trainable = False

```

```

# Creating a placeholder for the input of gan which would be a random vector of 50 dimension
gan_input = Input(shape =(latent_dim,))
# The image generated from the generator, which is then fed to discriminator
gen_img = generator(gan_input)
gan_output = discriminator(gen_img)

gan_model = Model(inputs = gan_input, outputs = gan_output)
gan_model.compile(loss = "binary_crossentropy", optimizer = keras.optimizers.Adam(learning_rate=0.0001))
return gan_model

gan_mod = gan(g, d)
gan_mod.summary()

```

Model: "functional_1"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 50)]	0

sequential (Sequential)	(None, 28, 28, 1)	541825

sequential_1 (Sequential)	(None, 1)	94049
=====		
Total params: 635,874		
Trainable params: 541,825		
Non-trainable params: 94,049		

```

def plot_images(epoch, generator, n_samples = 50, dim = (5, 10), figsize = (10, 10), latent_dim = latent_dim):
    """
    This function plots images generated using noise, it also saves the plots and the models for each epoch
    """
    # generating random noise
    noise = np.random.normal(0, 1, size = [n_samples, latent_dim])
    #generating images from the noise using generator
    gen_imgs = generator.predict(noise)
    gen_imgs = gen_imgs.reshape((n_samples, 28, 28))
    plt.figure(figsize = figsize)
    for i in range(n_samples):
        plt.subplot(dim[0], dim[1], i+1)
        plt.imshow(gen_imgs[i], cmap = 'gray_r', interpolation='nearest')
        plt.axis("off")
    plt.tight_layout()
    filename = 'gan_generated_image%d.png' %epoch
    plt.savefig(filename)
    filename = 'generator_model_%03d.h5' %epoch
    generator.save(filename)

```

```

def training_gan(epochs = 2, batch_size = 128):
    """
    This function trains the gan framework, by training discriminator and generator sequentially
    """

```

```

x_train, y_train, x_test, y_test = data_prepare()
g = generator()
d = discriminator()
gan_mod = gan(g, d)
batch_count = int(x_train.shape[0]/batch_size)
for e in range(1, epochs + 1):
    print("Epoch", e)
    for _ in range(batch_count):
        #generating random noise vector and passing it as input to generator for generating images
        noise = np.random.normal(0, 1, size = [batch_size, latent_dim])
        gen_imgs = g.predict(noise)
        # Real images taken randomly
        indexes = np.random.randint(0, x_train.shape[0], size = batch_size)
        real_imgs = x_train[indexes]
        #concatenating real and generated images and making labels for them
        training_imgs = np.concatenate([real_imgs, gen_imgs])
        y_batch = np.zeros(2*batch_size)
        y_batch[:batch_size] = 1
        d.trainable = True
        #Training the discriminator only
        d.train_on_batch(training_imgs, y_batch)

        #Tricking by labeling the noisy data as real data
        noise = np.random.normal(0, 1, size= [batch_size, latent_dim])
        y_gen = np.ones(batch_size)

        # Training generator by fixing discriminator
        d.trainable = False
        gan_mod.train_on_batch(noise, y_gen)
    if(e==1 or e%10 == 0):
        plot_images(e, g)

```

```

# call to train gan for 50 epochs with a batch size of 128 and print the resultant images
training_gan(50, 128)

```

3 9 9 0 3 0 8 6 5 7

7 6 8 4 4 0 6 8 1 5

4 6 2 0 5 5 4 4 0 2

0 1 7 3 3 5 4 6 4 9

0 5 8 2 3 9 1 9 6 0

6 0 4 6 6 4 8 7 2 4

1 5 6 5 2 7 1 2 3 0

4 7 7 2 2 1 0 8 9 8

9 7 5 1 1 0 6 9 0 3

9 9 4 4 C O 4 6 5 1